

# DataStax Developer Day

---



## Application Development



Source Code Inside

# Application Development



# Connectivity

DataStax Drivers

Cluster Builder

Contact Points

Configuration File



# DataStax Drivers

- OSS Driver Features

- CQL Support
- Sync / Async API
- Load Balancing Policies
- Retry Policies
- Reconnection Policies
- Connection Pooling
- SSL
- Compression
- Query Builder
- Object Mapper

- Enterprise Driver Features

- OSS Driver features, plus...
- DSE Advanced Security, Unified Authentication
- DSE Graph Fluent API
- DSE Geometric Types



- OSS Driver

```
<dependency>  
  <groupId>com.datastax.oss</groupId>  
  <artifactId>java-driver-core</artifactId>  
</dependency>
```

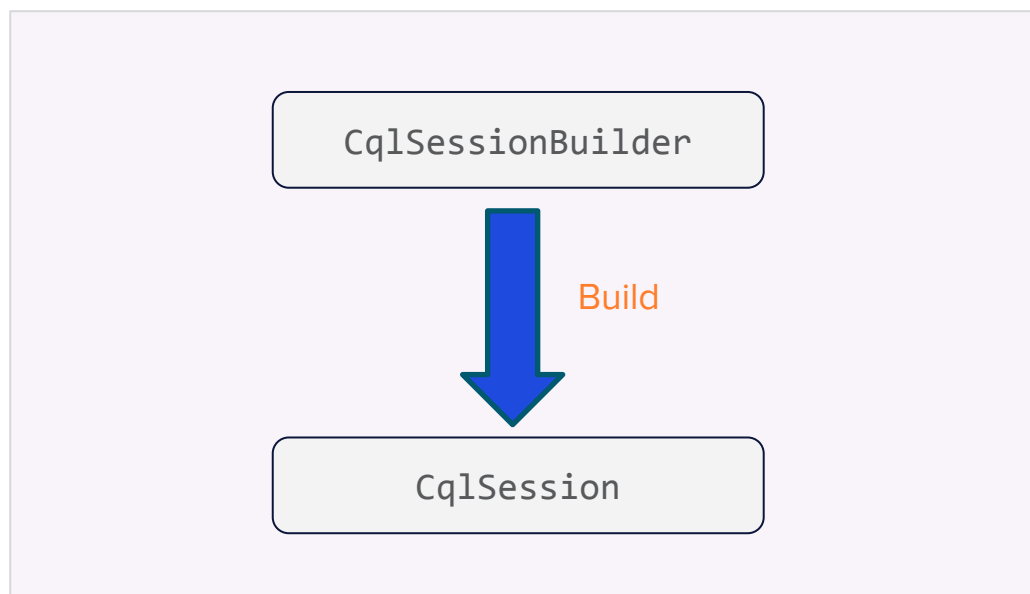
- DSE Driver

```
<<dependency>  
  <groupId>com.datastax.dse</groupId>  
  <artifactId>dse-java-driver-core</artifactId>  
</dependency>
```

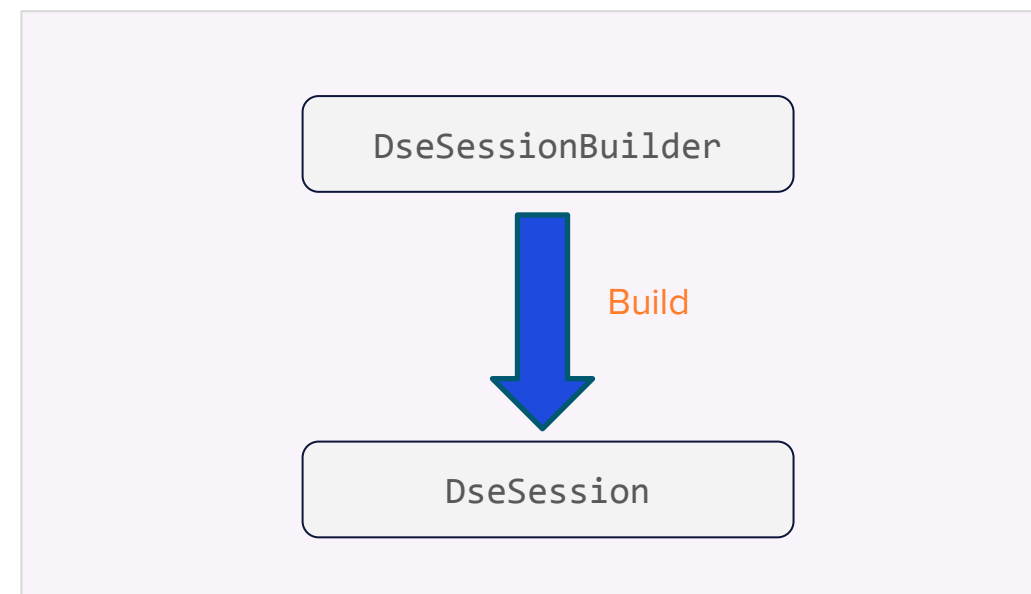


# Connectivity

- OSS Driver



- DSE Driver



*NB : “Cluster” concept from previous driver versions has been collapsed into “Session”*



# Builder ?





# Builder

// Explicit Settings


```
CqlSession cqlSession = CqlSession.builder()  
    .addContactPoint(new InetSocketAddress("127.0.0.1", 9042))  
    .withKeyspace("killrvideo")  
    .withLocalDatacenter("localDc")  
    .build();
```

// Delegate all configuration to file

```
CqlSession cqlSession = CqlSession.builder().build();
```

# Contact Points

- Only one necessary
- Unless that node is down
- More are good





# File-based Configuration

- Based on Typesafe Config
- Attributes are grouped into basic and advanced categories
- A reference file (*reference.conf*) provide default values embedded in the jar file. Can be override with key in application.conf.
- Driver searches `application.conf` in the classpath




```
datastax-java-driver {  
  basic {  
    request.timeout      = 5 seconds  
    request.consistency = LOCAL_QUORUM  
  }  
}
```



# Load-Balancing

- Used to create query plans for each statement executed
- Default policy is token-aware, round robin
- Requests are routed to nodes in the “local” data center only



application.conf

```
datastax-java-driver {  
  basic {  
    load-balancing-policy {  
  
      # The class of the policy.  
      class = DefaultLoadBalancingPolicy  
  
      # The datacenter that is considered "local"  
      # The default policy will only include nodes from  
      # this datacenter in its query plans.  
      local-datacenter = datacenter1  
  
      # A custom filter to include/exclude nodes  
      // filter.class=  
    }  
  }  
}
```



# Retry Policy


Determines when queries are retried on failure

- **DefaultRetryPolicy**
  - Default
  - Retries once onReadTimeout or onWriteTimeout
  - Enough replicas for your consistency level must be online
  - Only retries idempotent mutations



application.conf

```
datastax-java-driver {  
  
    # The policy that controls if the driver retries  
    # requests that have failed on one node.  
    advanced.retry-policy {  
  
        # The class of the policy  
        class = DefaultRetryPolicy  
  
    }  
}
```




# Reconnection Policy

Reconnects driver to a downed node

Two options:

- **ConstantReconnectionPolicy**
  - Check every N milliseconds
- **ExponentialReconnectionPolicy**
  - Increases every interval
  - Caps out at a max



application.conf

```
datastax-java-driver {  
  
    # Whether to schedule reconnection attempts  
    # if all contact points are unreachable at init  
    advanced.reconnect-on-init = false  
  
    advanced.reconnection-policy {  
  
        # The class of the policy  
        class = ExponentialReconnectionPolicy  
  
        # Parameters  
        base-delay = 1 second  
        max-delay = 60 seconds  
    }  
}
```



# Important to know about CqlSession

- **CqlSession** is a stateful object handling communications with each node
- **CqlSession** should be unique in the Application (*Singleton*)
- **CqlSession** should be closed at application shutdown (*shutdown hook*) in order to free opened TCP sockets (*stateful*)

```
@PreDestroy  
public void cleanup() {  
    if (null != cqlSession) {  
        cqlSession.close();  
    }  
}
```







# Time for an exercise!

---

## “Application Development” Notebook Exercise 1



# Application Development

DATASTAX<sup>®</sup>

## Building your queries

Simple Statements

Prepared Statements

Query Builder



# How to to execute queries ?

- First job of **CqlSession** is to execute queries using, well, execute method.

```
cqlSession.execute("SELECT * FROM killrvideo.users");
```



Statement



# SimpleStatement

```
Statement statement = ...
```

```
// (1) Explicit SimpleStatement Definition
```

```
SimpleStatement.newInstance("select * from t1 where c1 = 5");
```

```
// (2) Externalize Parameters (no name)
```

```
SimpleStatement.builder("select * from t1 where c1 = ?")  
    .addPositionalValue(5);
```

```
// (3) Externalize Parameters (name)
```

```
SimpleStatement.builder("select * from t1 where c1 = :myVal")  
    .addNamedValue("myVal", 5);
```

```
cqlSession.execute(statement);
```



# Prepared and Bound Statements

- Compiled once on each node automatically as needed
- Prepare each statement only once per application
- Use one of the many bind variations to create a BoundStatement

```
PreparedStatement ps = cqlSession.prepare("SELECT * from t1 where c1 = ?");  
  
BoundStatement bound = ps.bind(5);  
  
cqlSession.execute(bound);
```





# Query Builder



# Query Builder

- Fluent API for building CQL string queries programmatically
- Contains methods to build SELECT, UPDATE, INSERT and DELETE statements
- Generates a Statement as per the earlier techniques

## OSS Driver (current version 4.2.0)

```
<dependency>  
  <groupId>com.datastax.oss</groupId>  
  <artifactId>  
    java-driver-query-builder  
  </artifactId>  
</dependency>
```

## DSE Driver (current version 2.2.0)

```
<<dependency>  
  <groupId>com.datastax.dse</groupId>  
  <artifactId>  
    dse-java-driver-query-builder  
  </artifactId>  
</dependency>
```





# Query Builder

```
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.bindMarker;
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.deleteFrom;
import static com.datastax.oss.driver.api.querybuilder.QueryBuilder.selectFrom;
import static com.datastax.oss.driver.api.querybuilder.relation.Relation.column;
```

// Simple SELECT using QueryBuilder

```
Statement stmtSelect = selectFrom("killrvideo", "videos_by_users")
    .column("userid").column("commentid")
    .function("toTimestamp", Selector.column("commentid")).as("comment_timestamp")
    .where(column("userid").isEqualTo(bindMarker("userid")))
    .build();
```

// Simple DELETE using QueryBuilder

```
Statement stmtDelete = deleteFrom("killrvideo", "videos_by_users")
    .where(column("userid").isEqualTo(bindMarker("userid")))
    .build();
```



# Query Builder

- Can also use **QueryBuilder** to create **PreparedStatement**s and later execute at runtime
- Note use of **bindMarker()** to designate parameters that will be provided later

```
// Prepared QueryBuilder statements as any statement
PreparedStatement psStmt = cqlSession.prepare(
    deleteFrom("killrvideo", "videos_by_users")
        .where(column("userid").isEqualTo(bindMarker("userid")))
        .build());

// Binding
BoundStatement bsStmt = psStmt.bind("e7a8ac9f-c12d-415c-a526-4137815df573");


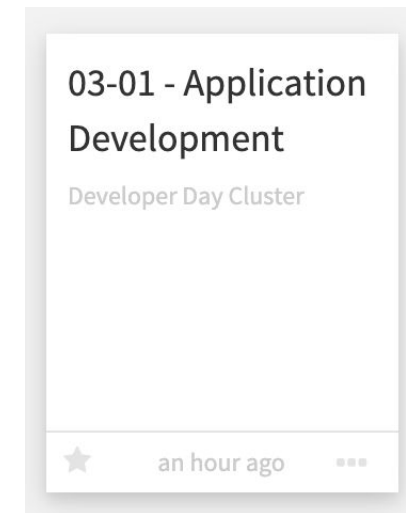
// Execute
cqlSession.execute(bsStmt);
```



# Time for an exercise!

---

## “Application Development” Notebook Exercises 2 and 3





# Application Development

DATASTAX<sup>®</sup>

## Executing Statements

Result Set

Parsing Rows

Batches

Profiles



# ResultSet

- ResultSet is the object returned for executing query. It contains **ROWS** (data) and **EXECUTION INFO**.
- ResultSet is **iterable** and as such you can navigate from row to row.
- Results are **always paged** for you (avoiding memory and response time issues)

```
ResultSet rs = cqlSession.execute(myStatement);

// Plumbery
ExecutionInfo info = rs.getExecutionInfo();
int executionTime = info.getQueryTrace().getDurationMicros();

// Data: NOT ALL DATA RETRIEVED IMMEDIATELY (only when needed .next())
Iterator<Row> iterRow = rs.iterator();
int itemsFirstCall = rs.getAvailableWithoutFetching();
```



# Parsing ResultSet

```
// We know there is a single row (eg: count)
Row singleRow = resultSet.one();

// We know there are not so many results we can get all (fetch all pages)
List<Row> allRows = resultSet.all();

// Browse iterable
for(Row myRow : resultSet.iterator()) {
    // .. Parsing rows
}

// Use Lambda
rs.forEach(row -> { row.getColumnDefinitions(); });

// Use for LWT
boolean isQueryExecuted = rs.wasApplied();
```



# Parsing Rows

```
// Sample row
Row row = resultSet.one();

// Check null before read
Boolean isUserNameNull = row.isNull("userName");

// Reading Values from row
String userName1 = row.get("username", String.class);
String userName2 = row.getString("username");
String userName3 = row.getString(CqlIdentifier.fromCql("username"));

// Tons of types available
row.getUuid("userid");
row.getBoolean("register");
row.getCqlDuration("elapsed");
...
```



# Paging

- ResultSet contains up to “[pageSize](#)” items. When browsing records you may hit this number that will trigger fetching next “pageSize” items.
- To fetch anything else that first page you must provide a [PagingState](#).

```
// Enforce few items per page (often = UI requirements)
myStatement = myStatement.setPageSize(10);
ResultSet page1 = cqlSession.execute(myStatement);

// Paging State
ByteBuffer pagingState = page1.getExecutionInfo().getPagingState();
myStatement = myStatement.setPageState(pagingState);

// Very same statement with pagingState provided
ResultSet page2 = cqlSession.execute(myStatement);
```





# Batches – What you need to know

- Batches about **data integrity** between tables
- Not Atomic & not used for mass query optimization
- Used to keep denormalized data in sync
- There is no guarantee that a batch will complete all operations.
- There are still edge cases where things can fail out.
- There is no rollback if something fails
- This is where upserts come into play as you can simply re-fire the batch.



# Batch Example

```
// Sample statements (insert same data in multiple tables)
Statement stmt1 = SimpleStatement
    .builder("INSERT INTO users_by_group(groupid,userid) values(?,?)")
    .addPositionalValue(groupname, username);
Statement stmt2 = SimpleStatement
    .builder("INSERT INTO groups_by_user(userid,groupid) values(?,?)")
    .addPositionalValue(username, groupname);

// Group as a Batch
BatchStatement batchStmt = BatchStatement
    .builder(DefaultBatchType.LOGGED)
    .addStatement(stmt1).addStatement(stmt2).build();

// Execute
cqlSession.execute(batchStmt);
```



# Profiles

Override parameters for dedicated request

```
SimpleStatement
    .newInstance("select...")
    .setPageSize(10)
    // here is the magic ☐
    .setExecutionProfileName("dse_search");
```



application.conf

```
datastax-java-driver {
  profiles{



    # DSE Search type of queries
    dse_search {
      basic {
        request.consistency = LOCAL_ONE
        request.timeout      = 5 seconds
      }
    } # The class of the policy
    fast_query {
      basic.request.consistency = ONE
      basic.request.timeout     = 1 second
    }
  }
}
```



# Time for an exercise!

---

## “Application Development” Notebook Exercise 4



# Application Development

DATASTAX<sup>®</sup>

# Object Mapping

Entity

Dao

Mapper

Query Provider



# Object Mapper

- WHAT ?
  - Abstracts details of mapping Java attributes to/from CQL types and UDTs
  - Packaged separately from the driver – pom.xml update required
  - This slide shows the runtime dependency – will show compile-time shortly
- HOW ?
  - Some annotation processors will GENERATE Mapper, Dao, and Entity implementations for you
  - At each update in the files, the IDE (eclipse, intelliJ) will use annotation processor
  - Compiler plugin must be updated to define the annotation processor



# Object Mapper

## OSS Driver

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.oss</groupId>
      <artifactId>java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
  </configuration>
</plugin>
```

## DSE Driver

```
<dependency>
  <groupId>com.datastax.dse</groupId>
  <artifactId>dse-java-driver-mapper-runtime</artifactId>
</dependency>

<!-- X -->

<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <release>11</release>
  <annotationProcessorPaths>
    <path>
      <groupId>com.datastax.dse</groupId>
      <artifactId>dse-java-driver-mapper-processor</artifactId>
    </path>
  </annotationProcessorPaths>
  </configuration>
</plugin>
```





# Annotate Entities

```
@Entity
@CqlName("user_v")
public class UserVideo {

    @PartitionKey
    @CqlName("userid")
    private UUID userid;

    @ClusteringColumn(1)
    @CqlName("added")
    private UUID videoid;

}
```

← TABLE NAME, KEYSPACE

← PARTITION KEY COLUMNS

← CLUSTERING COLUMNS



# Annotate DAO Interface (1/2)

```
@Dao
public interface VideoDao {

    @Select
    Optional<UserVideo> findUserById(UUID userid);

    @Query("SELECT * FROM ${tableId}")
    PagingIterable<UserVideo> findAll();

    @Select(customWhereClause = "videoid = : videoid")
    PagingIterable<UserVideo>
    findUserByVideoId(@CqlName("videoid") UUID vid);
}
```



## Annotate DAO Interface (2/2)

```
// Save a bean
@Insert
void save(UserVideo userVideo);

// Userid id is PK
@Delete
void delete(UUID userid);

// Custom implementations
@QueryProvider(
    providerClass = MySampleQueryProvider.class,
    entityHelpers = { UserVideo.class })
String doSomething(String abc);
```



# Annotate Mapper Interface

```
@Mapper
public interface MyApplicationMapper {

    @DaoFactory
    VideoDao videoDao(@DaoKeyspace CqlIdentifier keyspace);

}
```



# Sample QueryProvider

```
public class MySampleQueryProvider {  
  
    // Constructor, getting session  
    public MySampleQueryProvider(  
        MapperContext context,  
        EntityHelper<UserVideo > helperUser) {}  
  
    // Custom implementation method  
    public String doSomething(String abc) {  
    }  
}
```



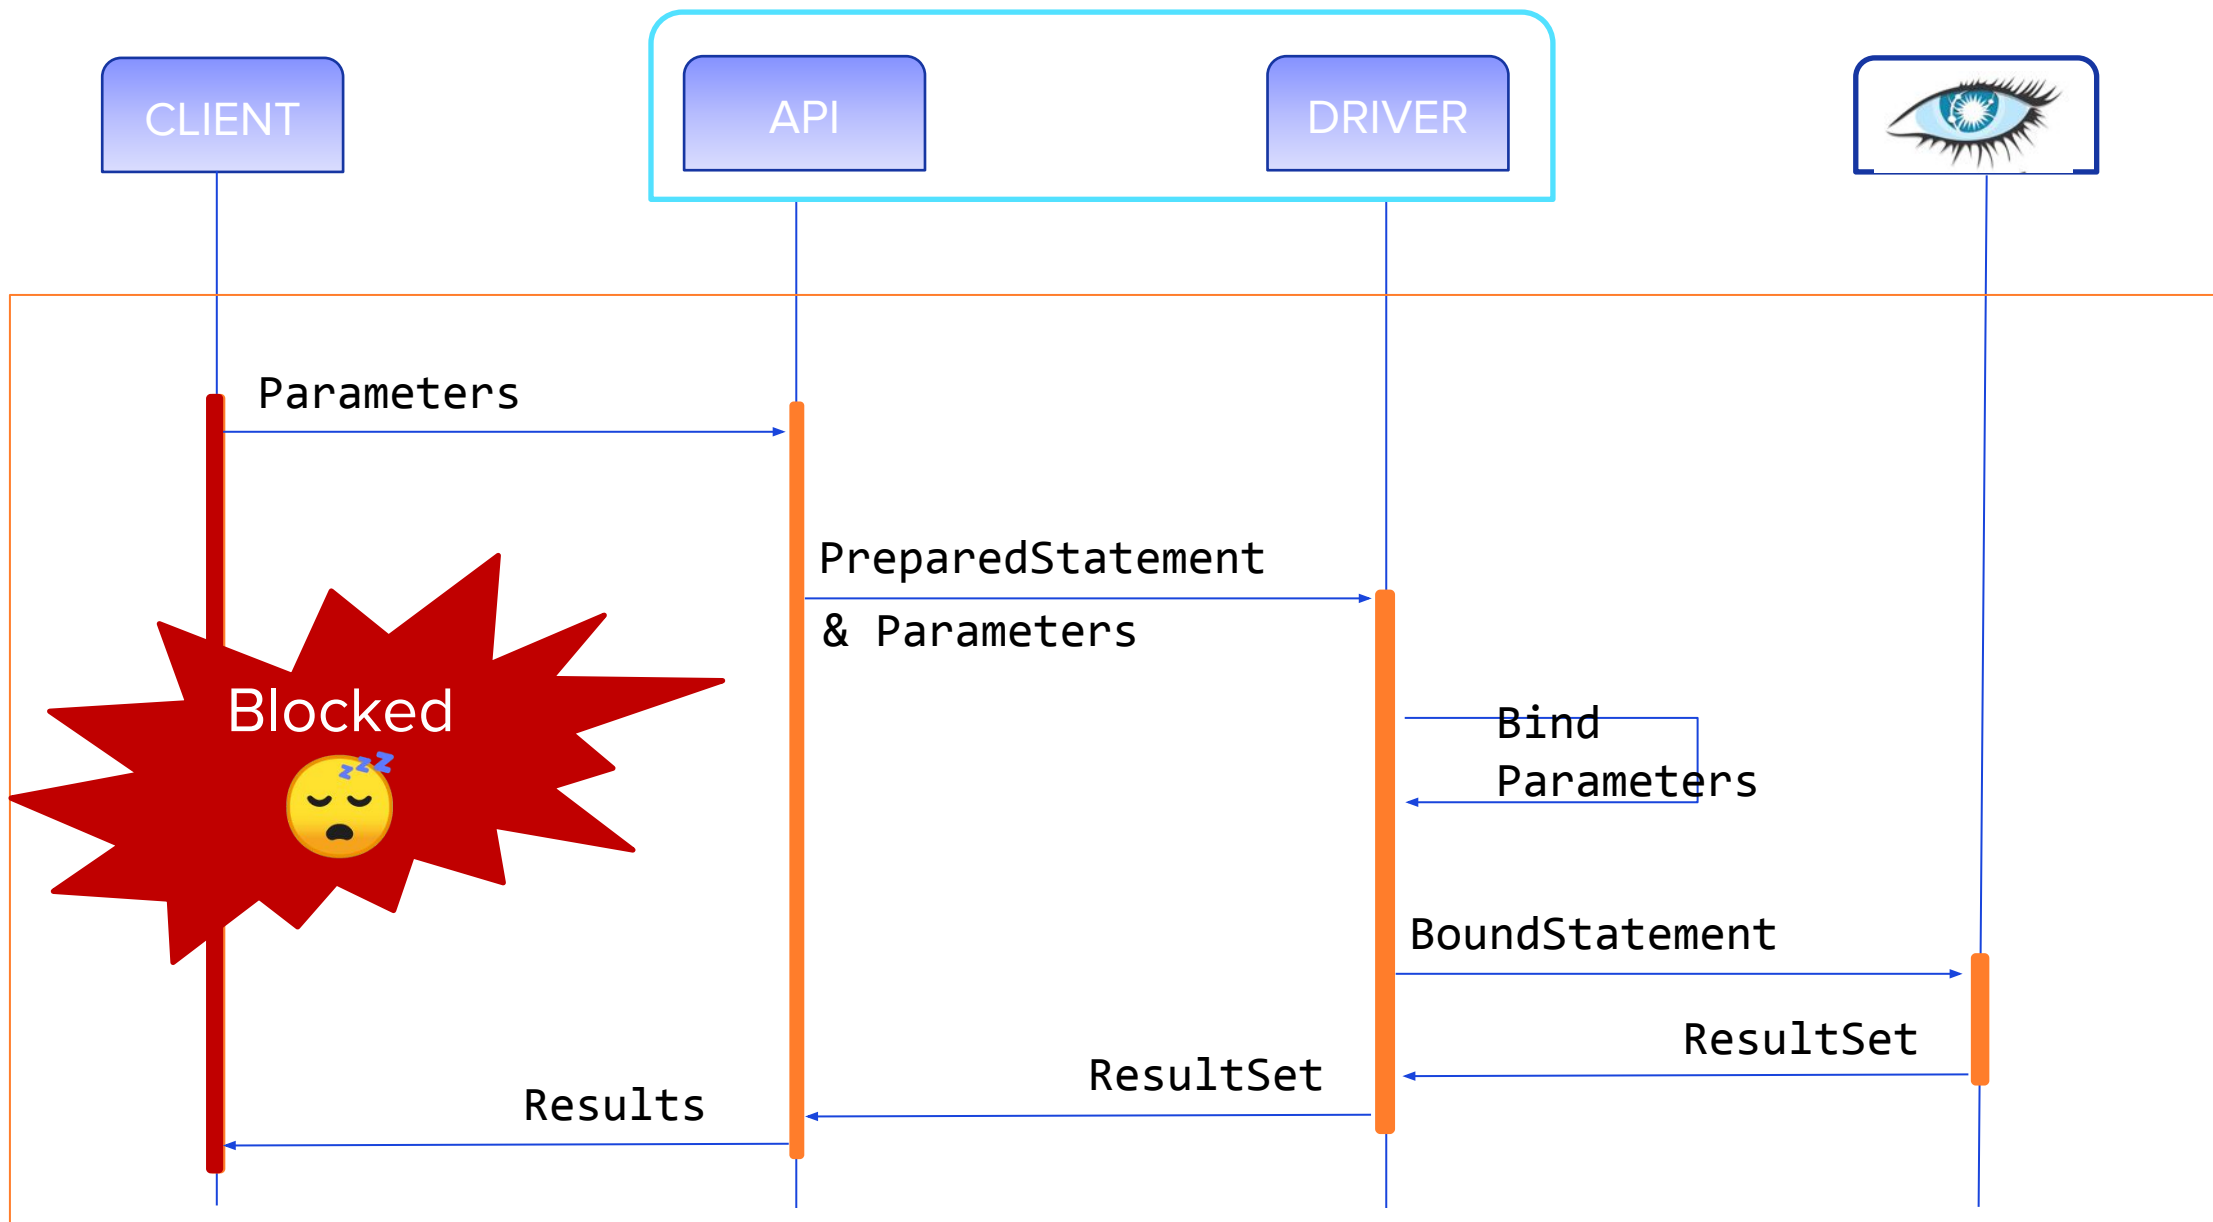
# Application Development



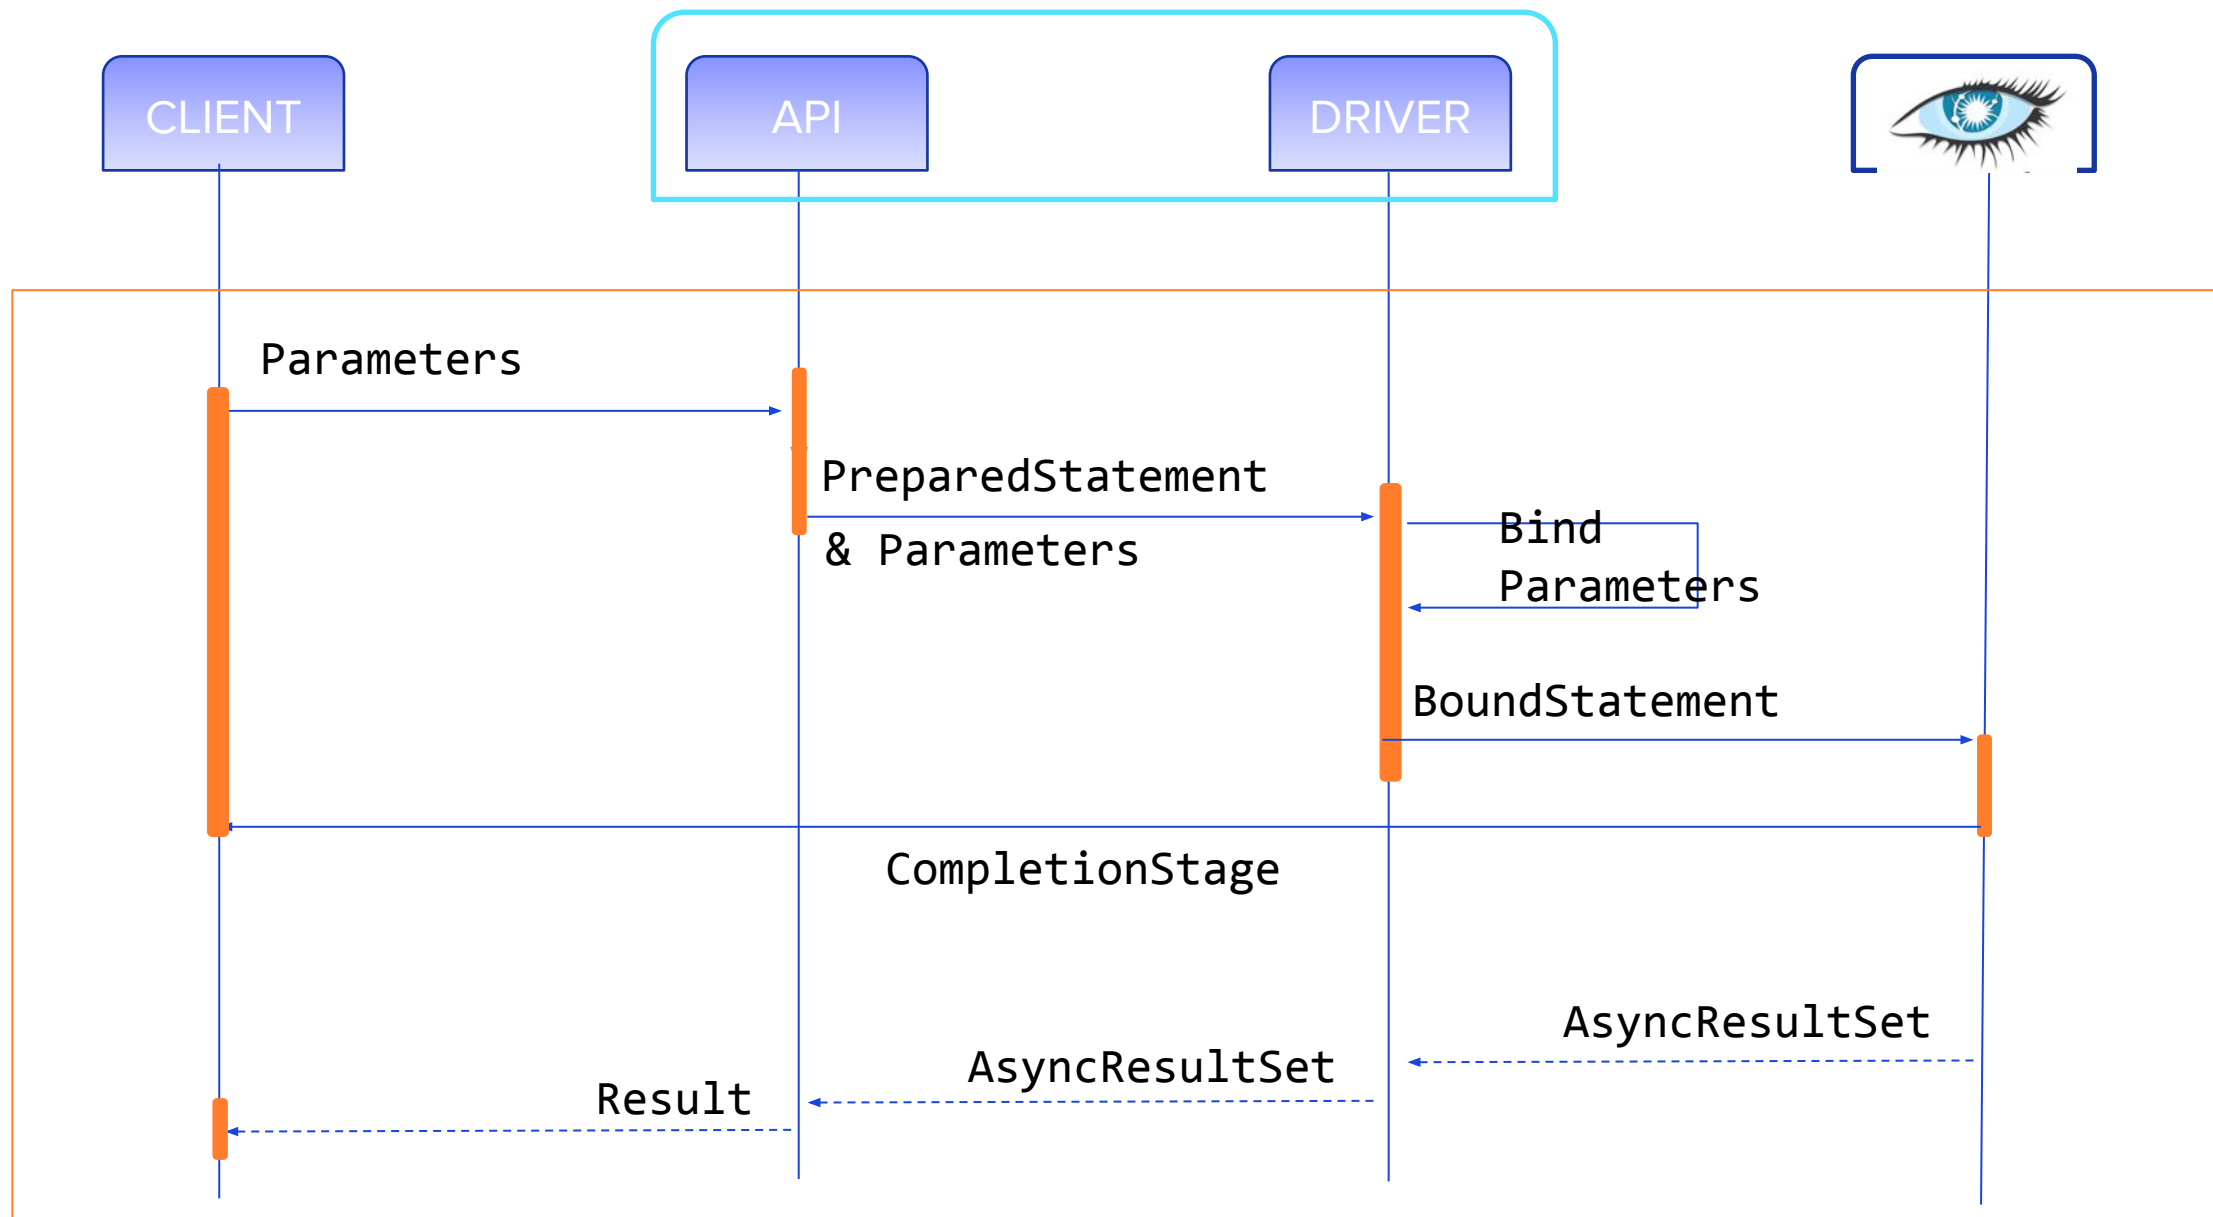
## Asynchronous APIs



Tell more about this section here







# Asynchronous Queries

```
// From Synchronous
ResultSet resSync = cqlSession.execute(myStatement);

// to Asynchronous
CompletionStage<AsyncResultSet> resAsync =
    cqlSession.executeAsync(myStatement);

resAsync.thenApply(AsyncPagingIterable::one)
        .thenApply(Optional::ofNullable)
        .thenApply(optional -> optional.map(rowMapper))
        .then..
```



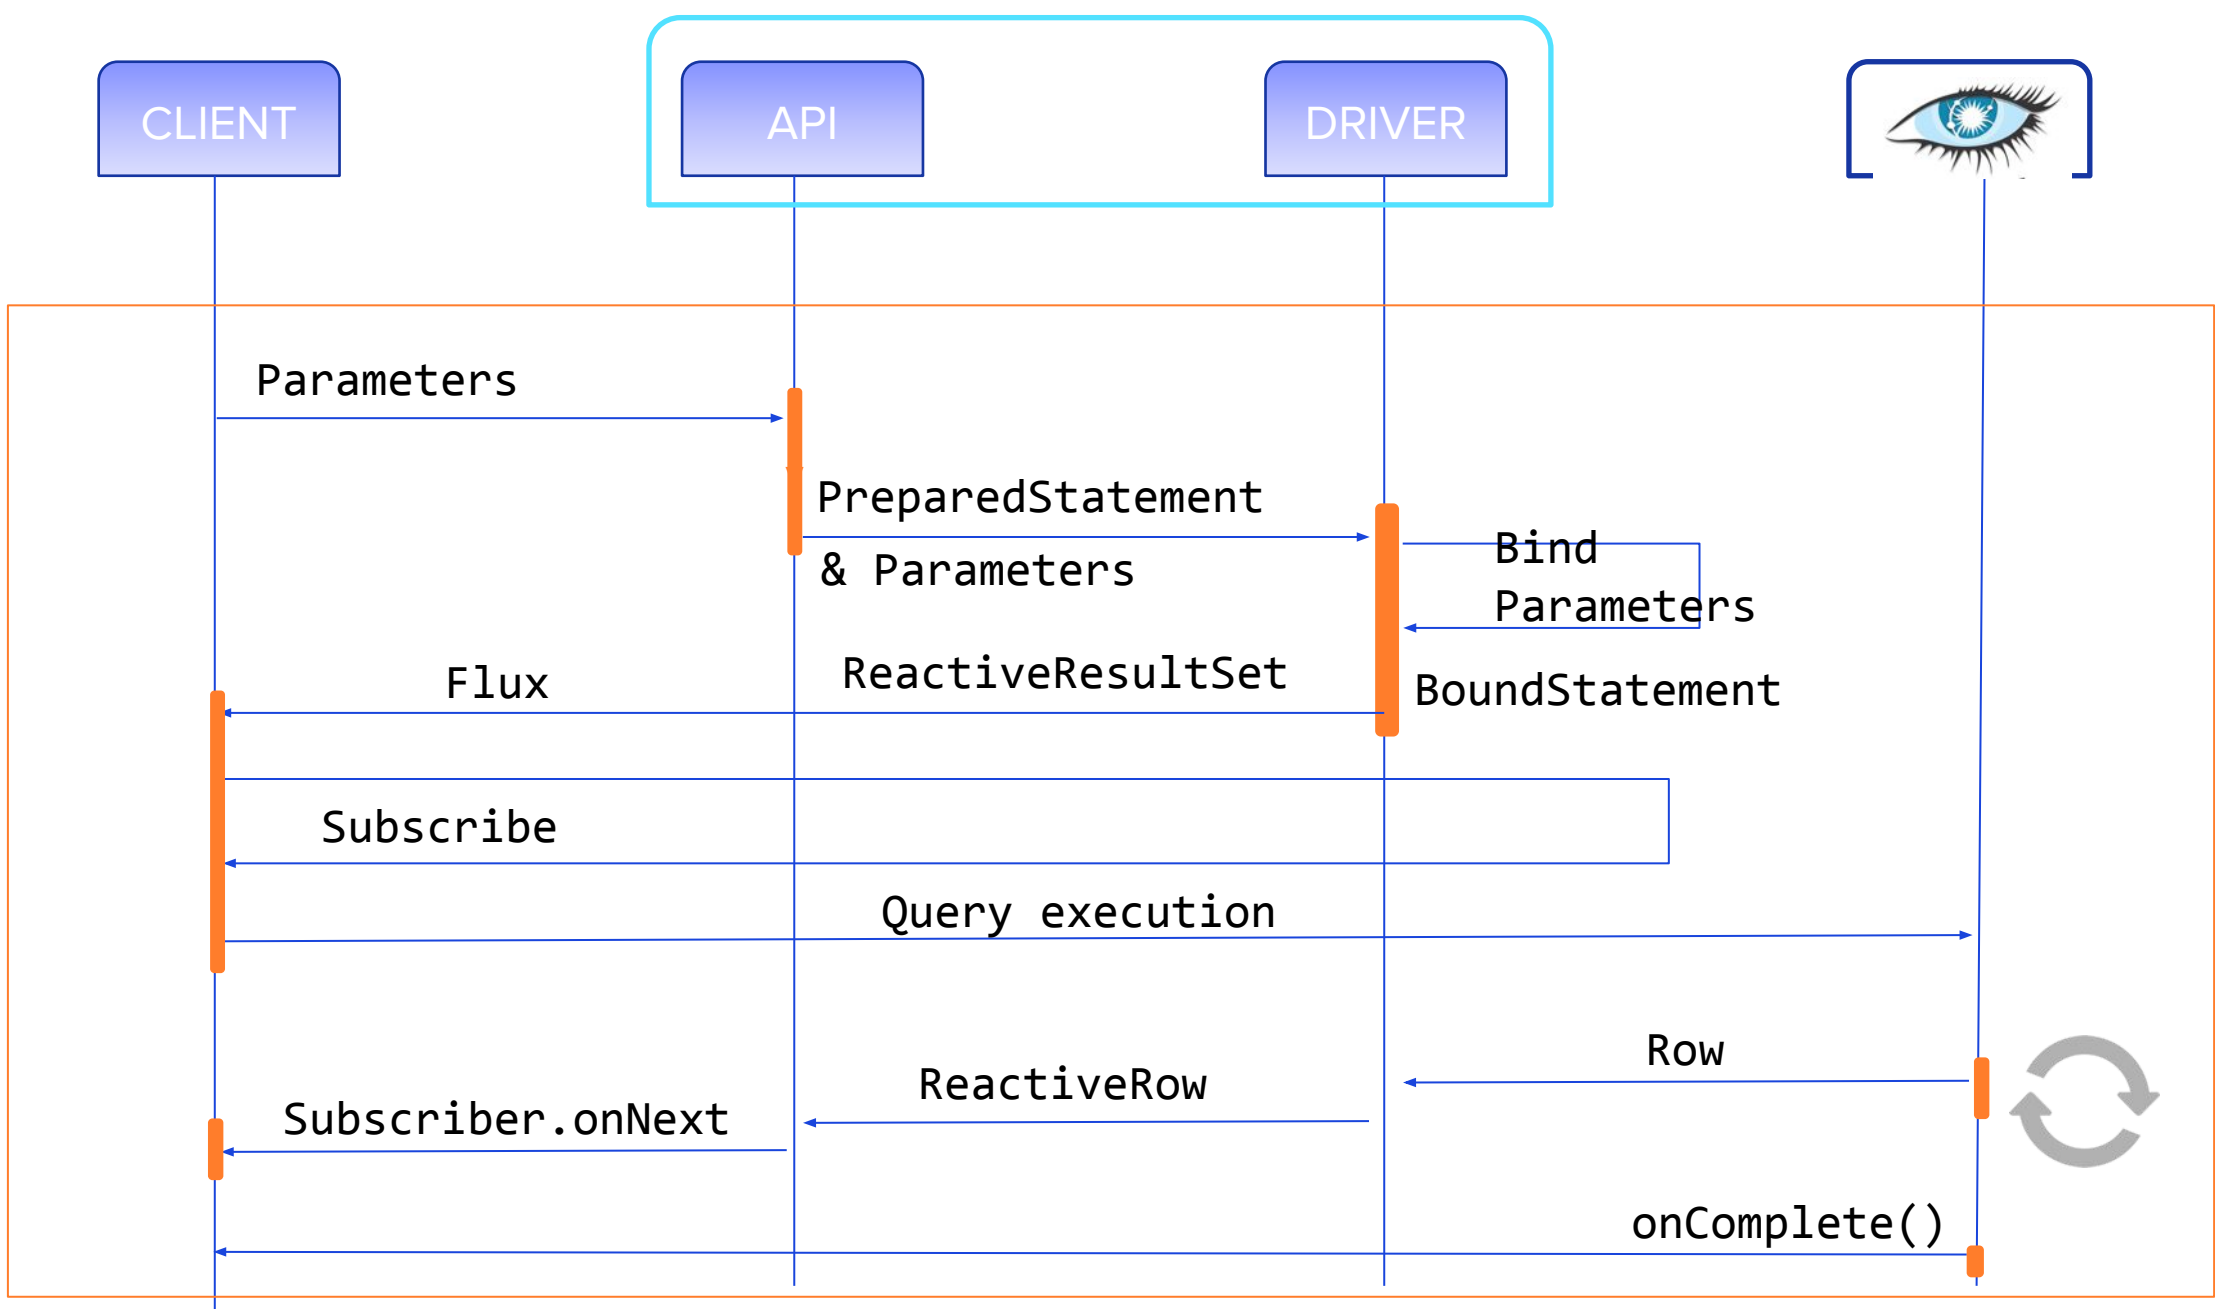
# Application Development



## Reactive APIs



Tell more about this section here



# Reactive Queries

```
private final Function<Row, MyBean> rowMapper = ..;  
  
// Execute in reactive way  
ReactiveResultSet rs = session.executeReactive(myStatement);  
  
// Return Flux for lists  
Flux<ReactiveRow> flux = Flux.from(rs);  
Flux<MyBean> res1= flux.skip(offset).take(limit).map(rowMapper);  
  
// Return Mono for single bean  
Mono<MyBean> res2 = Mono.fromDirect(rs).map(rowMapper);
```



# Application Development

DATASTAX<sup>®</sup>

## Spring Boot Starter



Tell more about this section here

# Spring Boot Starter

Define all configuration in `application.yaml` file

```
<dependency>
  <groupId>com.datastax.oss</groupId>
  <artifactId>
    java-driver-spring-boot-starter
  </artifactId>
  <version>1.0.0.20190903-LABS</version>
</dependency>
```



`application.yaml` \\_0\_/

```
datastax-java-driver:
  basic.contact-points:
    - 127.0.0.1:9042
  basic.session-keyspace: test
  basic.load-balancing-policy:
    local-datacenter: datacenter1
```

<https://github.com/datastax/labs/tree/master/spring-boot-starter/20190903>



@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>





# Application Development

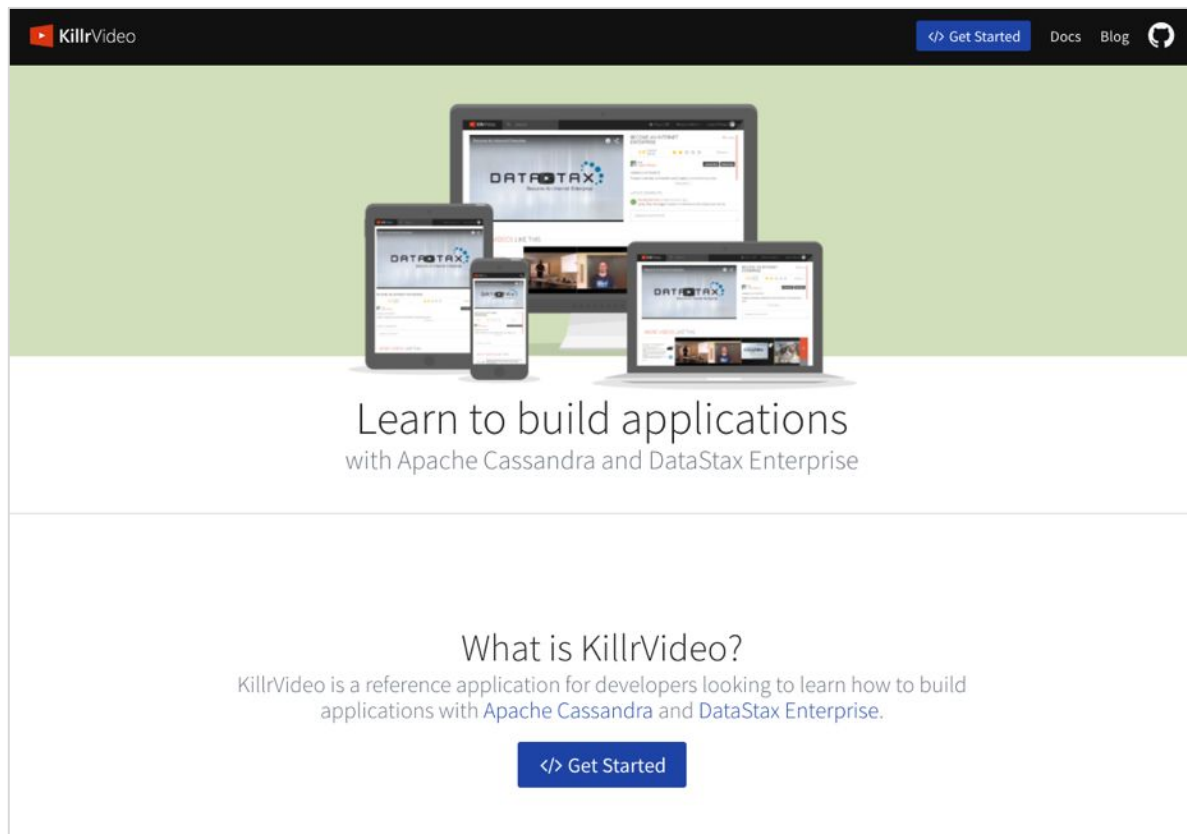


## Wrapping up

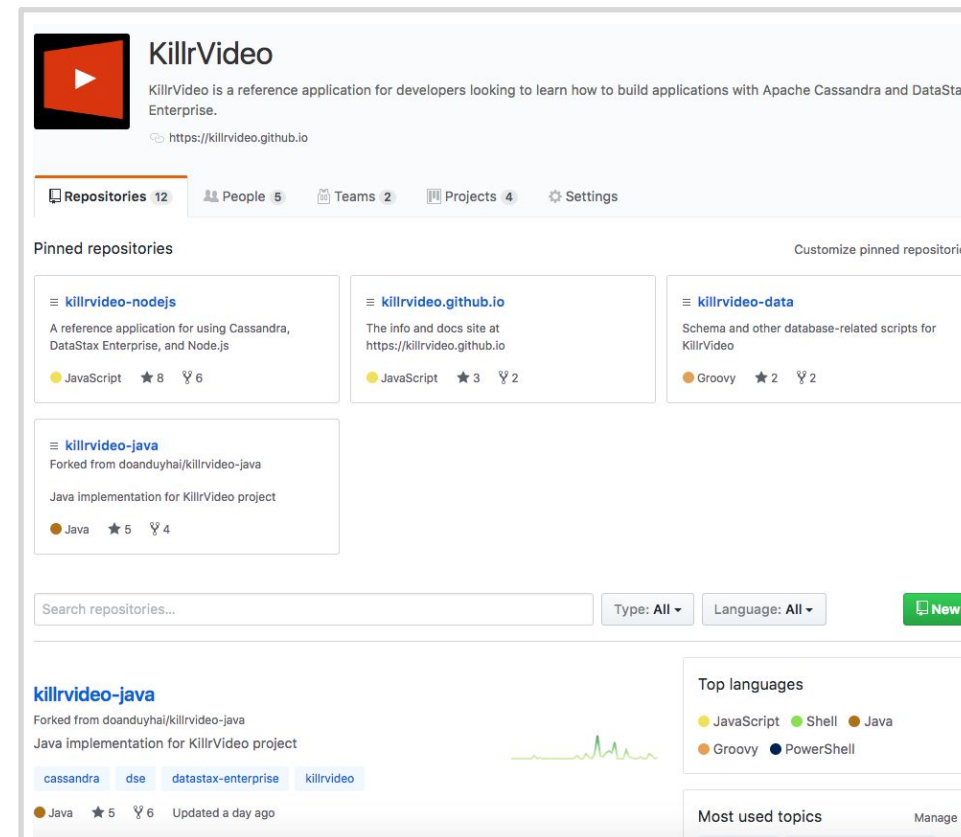


Tell more about this section here

# KillrVideo



<https://killrvideo.github.io>



<https://github.com/KillrVideo>

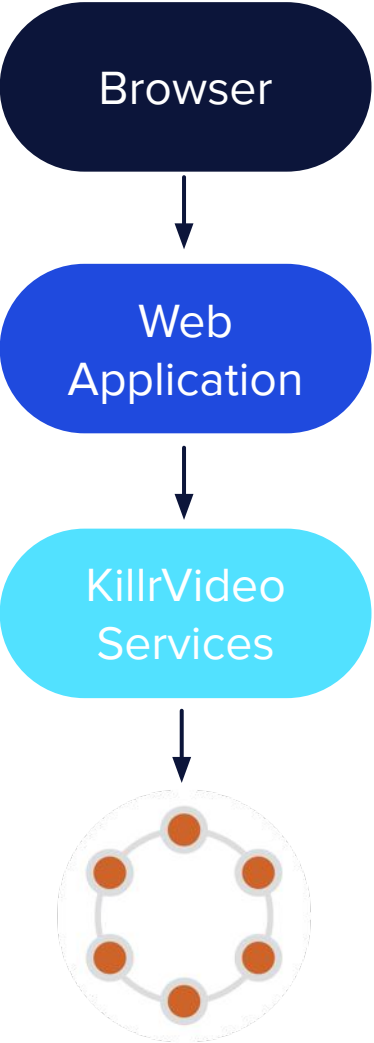


@DataStaxDevs #DataStaxDeveloperDay

<https://community.datastax.com>



# KillrVideo Architecture



## Technology Choices

- Node.js
- Falcor
- Java / C# / Node.js / Python
- GRPC
- DataStax Drivers
- DataStax Enterprise including Apache Cassandra & Spark, Graph

## Deployment

- Download and run locally via Docker
- Deployed in AWS using DataStax Managed Cloud: <http://killrvideo.com/>





# Thank You

