


# DataStax Developer Day

---

## Cassandra Data Modeling



DATASTAX®



# Cassandra Data Modeling

---

DATASTAX<sup>®</sup>

**SQL (relational) vs.  
CQL (Cassandra)**



# Structuring Your Database

- Normalization: To reduce data redundancy and increase data integrity.
- Denormalization: Must be done in read heavy workloads to increase performance



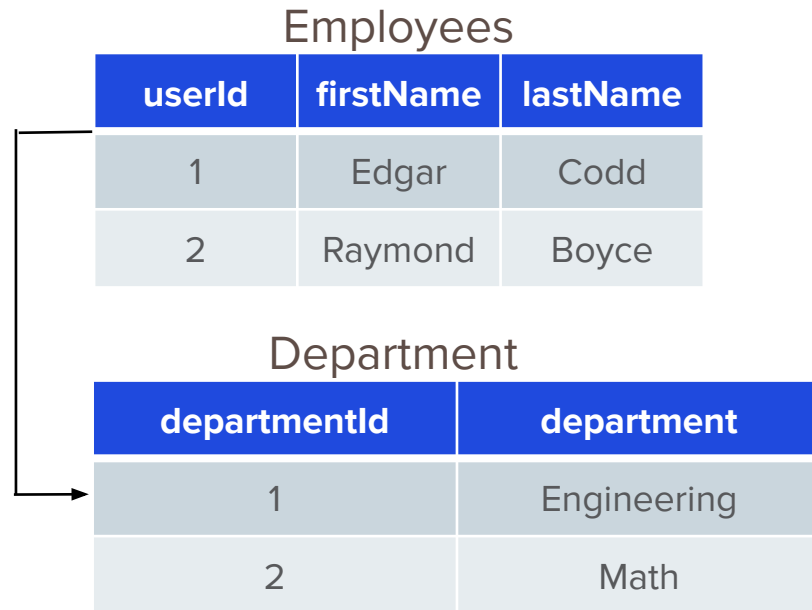
# Normalization

- Structuring a relational database
- Normal forms (3NF max)
- Why?
  - Reduce data redundancy
  - Increase data integrity.



# Relational Data Models

- Multiple normal forms
  - most do not go beyond 3NF
- Foreign Keys
- Joins



# Relational Modeling

- Create entity table
- Add constraints
- Index fields
- Foreign Key relationships



```
CREATE TABLE users (  
  id      number(12) NOT NULL ,  
  firstname  nvarchar2(25) NOT NULL ,  
  lastname  nvarchar2(25) NOT NULL,  
  email     nvarchar2(50) NOT NULL,  
  password  nvarchar2(255) NOT NULL,  
  created_date timestamp(6),  
  PRIMARY KEY (id),  
  CONSTRAINT email_uq UNIQUE (email)  
);  
  
-- Users by email address index  
CREATE INDEX idx_users_email ON users (email);
```

```
CREATE TABLE videos (  
  id number(12),  
  userid number(12) NOT NULL,  
  name nvarchar2(255),  
  description nvarchar2(500),  
  location nvarchar2(255),  
  location_type int,  
  added_date timestamp,  
  CONSTRAINT users_userid_fk  
    FOREIGN KEY (userid)  
    REFERENCES users (id) ON DELETE CASCADE,  
  PRIMARY KEY (id)  
);
```





# Relational Modeling



Employees

userId	firstName	lastName
1	Edgar	Codd
2	Raymond	Boyce


Department

departmentId	department
1	Engineering
2	Math






# Cassandra Modeling



id	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math



# Denormalization - Why?

- Improve read performance of a database
- Reduce write performance
  - Adding redundant copies of data



# CQL vs SQL

- No joins
- Limited aggregations

```
SELECT e.First, e.Last, d.Dept
FROM Department d, Employees e
WHERE 'Codd' = e.Last
AND e.deptId = d.id
```

Employees

userId	firstName	lastName
1	Edgar	Codd
2	Raymond	Boyce

Department

departmentId	department
1	Engineering
2	Math



# Applying Denormalization

- Combine table columns into a single view
- Eliminate the need for joins
- Queries are concise and easy to understand

Employees

id	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math

```
SELECT First, Last, Dept  
FROM employees  
WHERE id = '1'
```




# Denormalization in Apache Cassandra

- Denormalization of tables in Apache Cassandra is absolutely critical.
- The biggest take away is to think about your queries first.
- There are no JOINS in Apache Cassandra.




# Queries in Relational vs NoSQL Databases

- In a relational database, one query can access and join data from multiple tables



- In Apache Cassandra, you cannot join data, queries can only access data from one table

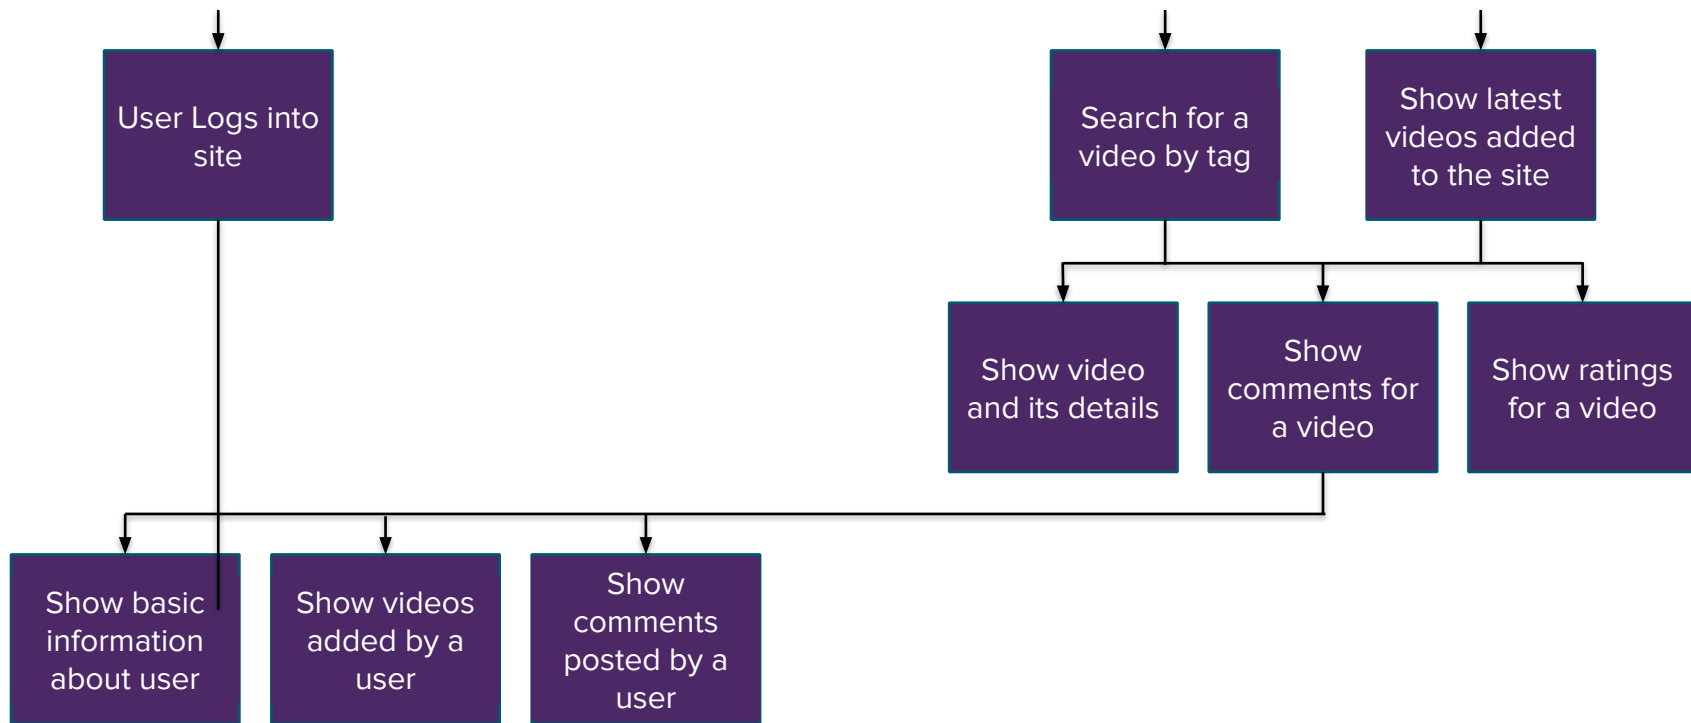


# Modeling Queries

- What are your application's workflows?
- Knowing your queries in advance is CRITICAL
- Different from RDBMS because I can't just JOIN or create a new indexes to support new queries
- One table per one query



# Some Application Workflows in KillrVideo





# Some Queries in KillrVideo to Support Workflows

## Users

User Logs into  
site

Find user by email  
address

Show basic  
information  
about user

Find user by id

## Comments

Show  
comments for  
a video

Find comments by  
video (latest first)

Show  
comments  
posted by a  
user

Find comments by user  
(latest first)

## Ratings

Show ratings  
for a video

Find ratings by video



# Cassandra Data Modeling

---

## Denormalization Mind Shift

DATASTAX<sup>®</sup>




# Cassandra Data Modeling Principles

- Design tables around queries
- Use partition key column(s) to group data you would like to be able to get in a single query
- Use clustering columns to guarantee unique rows and control sort order
- Use additional columns to provide the details you need
  - Denormalization - including data that might have been joined from elsewhere in a relational model



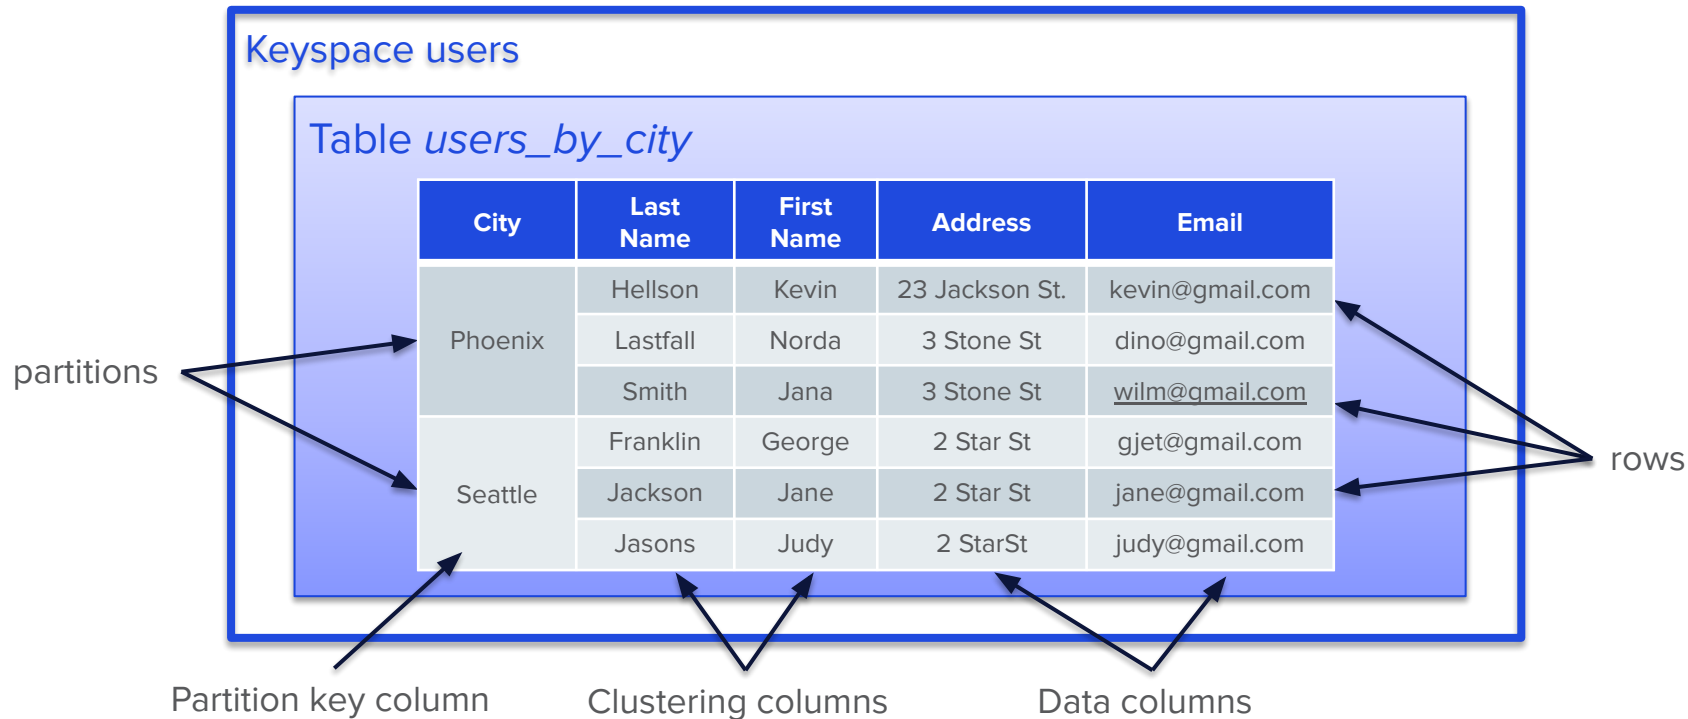
# Cassandra Structure - Partition



- Tabular data model, with one twist
- *Keyspaces* contain *tables*
- *Tables* are organized in *rows* and *columns*
- Groups of related rows called *partitions* are stored together on the same node (or nodes)
- Each row contains a *partition key*
  - One or more columns that are hashed to determine which node(s) store that data



# Example Data – Users organized by city



# Tables Hold Many Partitions

City	Last Name	First Name	Address	Email
Phoenix	Hellson	Kevin	23 Jackson St.	kevin@gmail.com
	Lastfall	Norda	3 Stone St	dino@gmail.com
	Smith	Jana	3 Stone St	wilm@gmail.com

Table *users\_by\_city*



# Tables Hold Many Partitions

City	Last Name	First Name	Address	Email
Seattle	Franklin	George	2 Star St	gjet@gmail.com
	Jackson	Jane	2 Star St	jane@gmail.com
	Jasons	Judy	2 StarSt	judy@gmail.com

Table *users\_by\_city*

City	Last Name	First Name	Address	Email
Phoenix	----	----	----	----
	----	----	----	----
	----	----	----	----



# Tables Hold Many Partitions

City	Last Name	First Name	Address	Email
Charlotte	Azrael	Chris	5 Blue St	chris@gmail.com
	Stilson	Brainy	7 Azure Ln	brain@gmail.com
	Smith	Cristina	4 Teal Cir	clu@gmail.com
	Sage	Grant	9 Royal St	grant@gmail.com
	Seterson	Peter	2 Navy Ct	peter@gmail.com

Table *users\_by\_city*

City	Last Name	First Name	Address	Email
Phoenix	----	----	----	----
	----	----	----	----
	----	----	----	----
Seattle	----	----	----	----
	----	----	----	----
	----	----	----	----





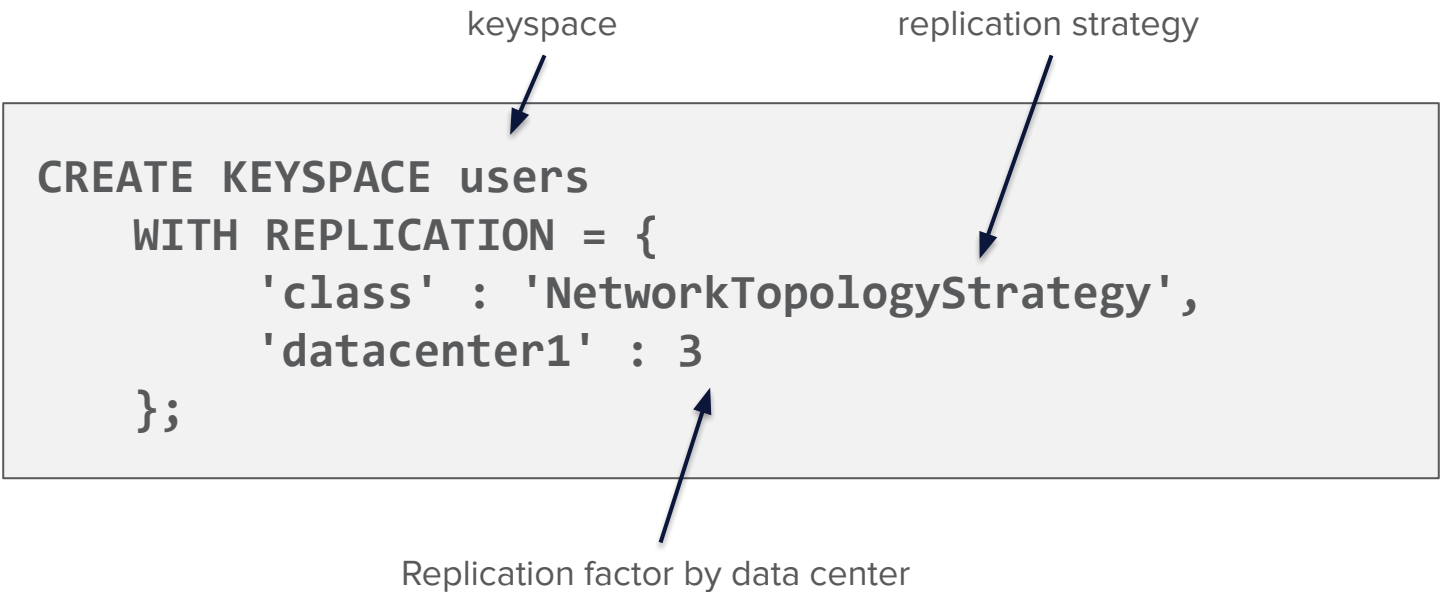
# Tables Hold Many Partitions

Table *users\_by\_city*

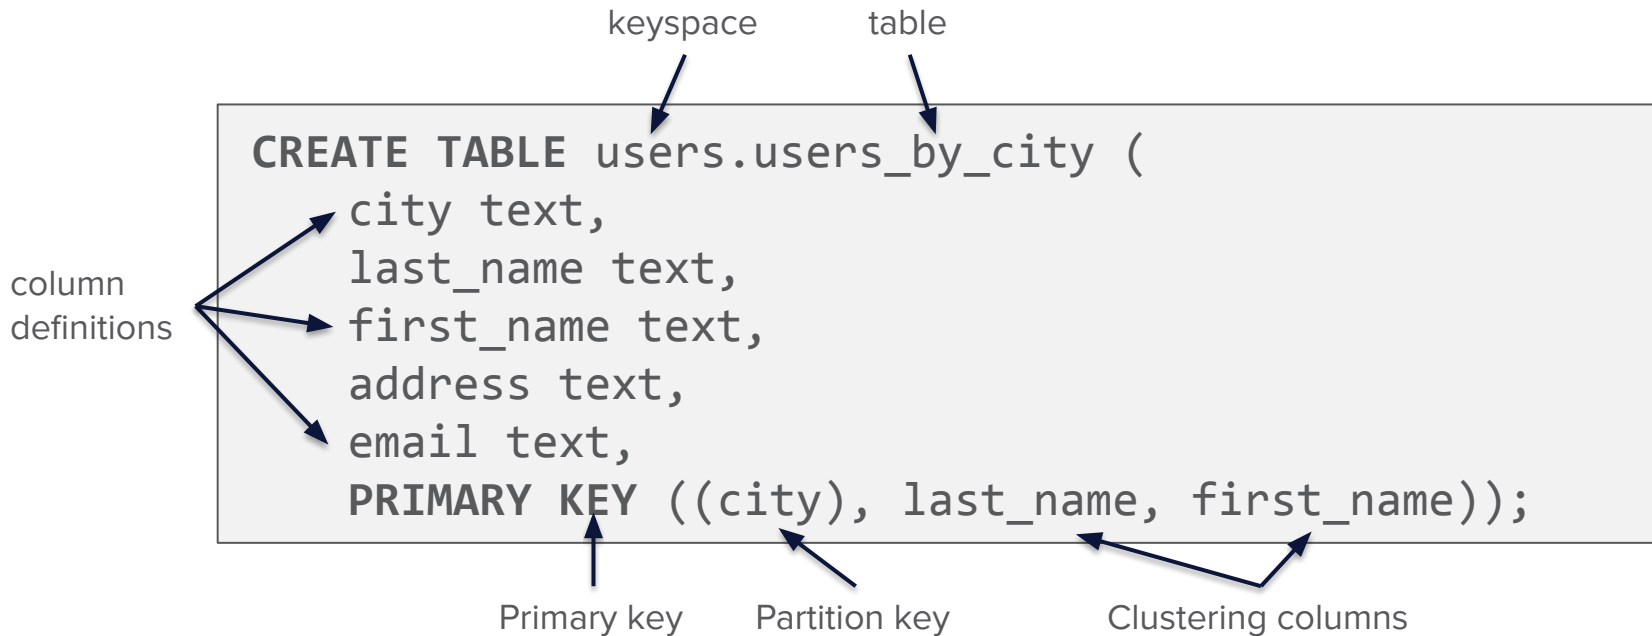
City	Last Name	First Name	Address	Email
Phoenix	----	----	----	----
	----	----	----	----
	----	----	----	----
Seattle	----	----	----	----
	----	----	----	----
	----	----	----	----
Charlotte	----	----	----	----
	----	----	----	----
	----	----	----	----
	----	----	----	----
	----	----	----	----



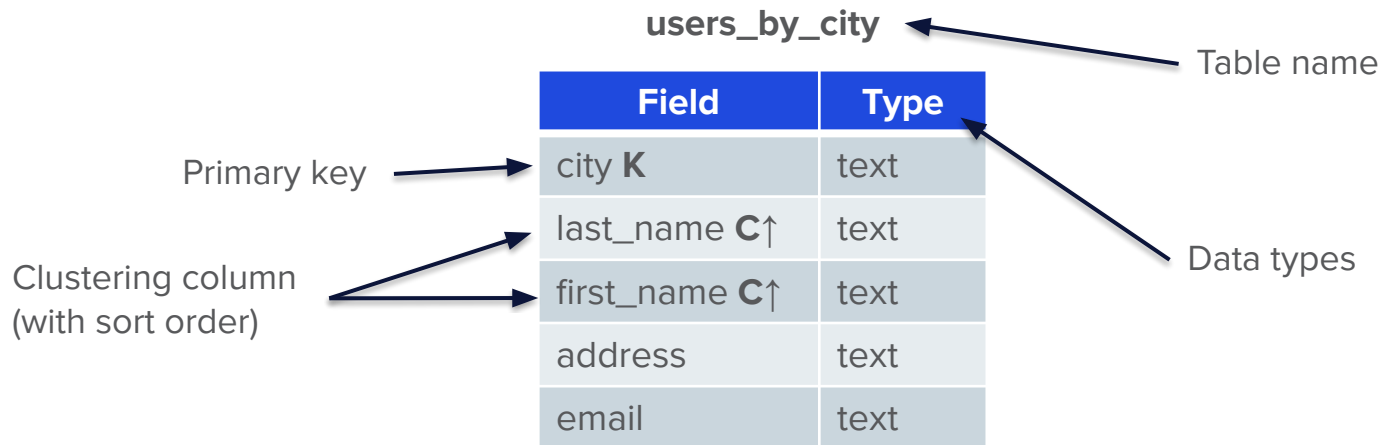
# Creating a Keyspace in CQL



# Creating a Table in CQL




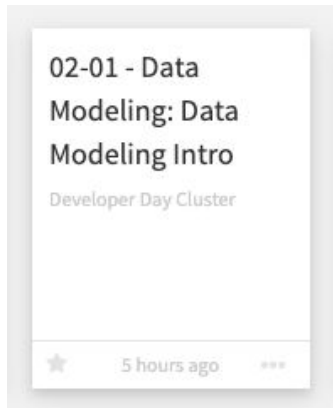
# Cassandra Table Design Notation - Chebotko Diagram



# Time for an exercise!

---

## “Data Modeling Intro” Notebook



# Data Modeling – Key Concepts

- Keyspace – contains tables
- Table – contains partitions
- Row – has a primary key and data columns
- Partition – basic unit of storage/retrieval
  - Identified by partition key embedded within primary key
  - Contains one or more rows
- Primary key – intra-table row identifier
  - Consists of partition key and clustering columns
  - Partition key – partition identifier, hashes to partition token
  - Clustering column – intra-partition key for sorting rows within partition



# Welcome to Cassandra-Land

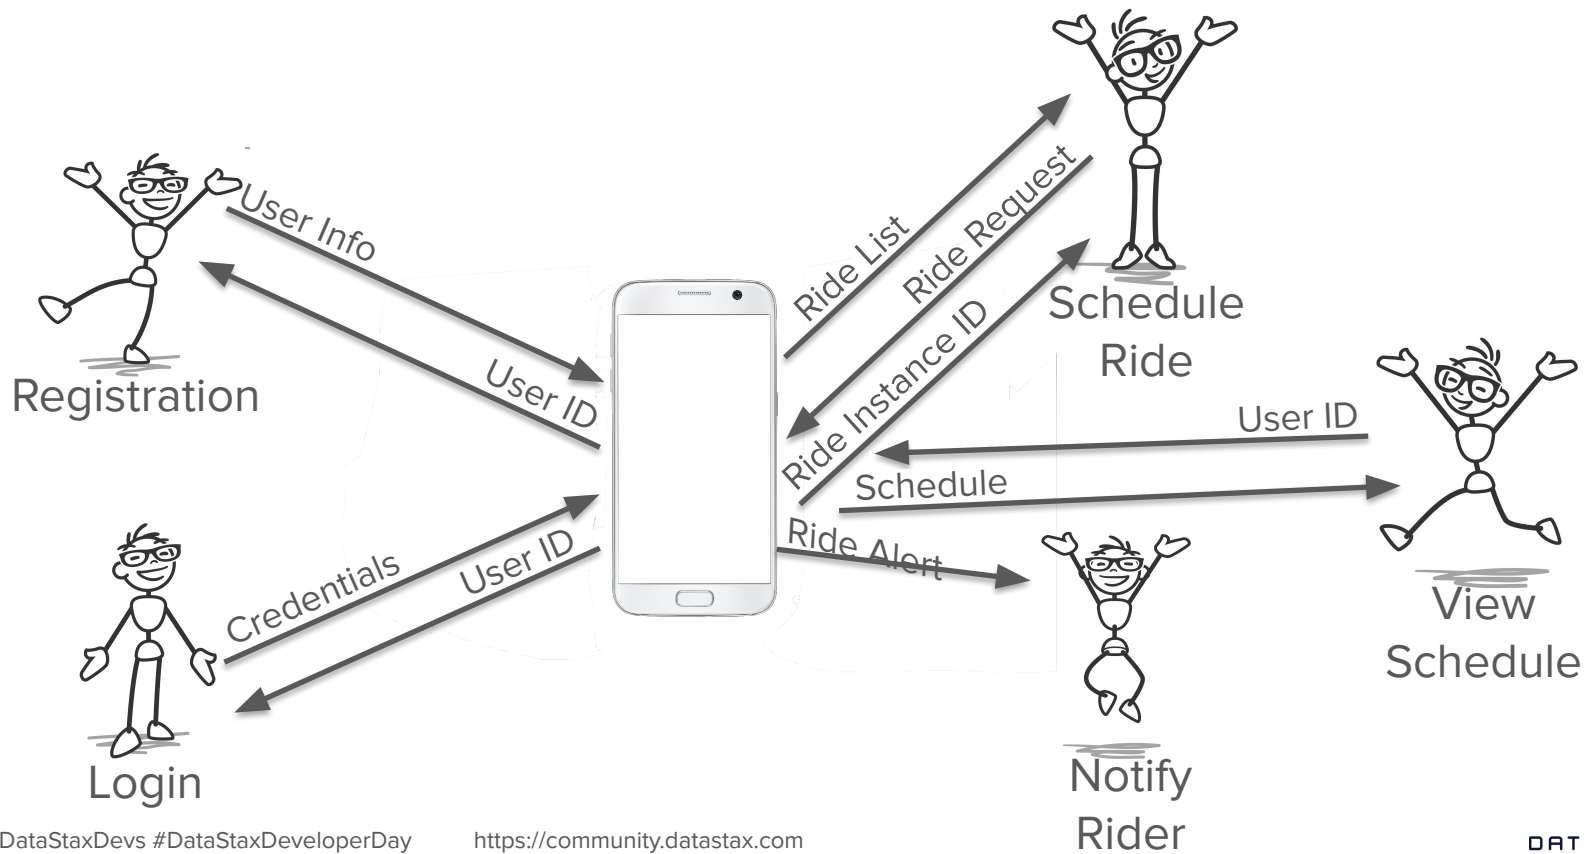
The Theme Park Where You Can Find...

- Distributed & Fault-Tolerant Rides
- Amazing Throughput
- And Fast Response Times

But We Need an App!



# Cassandra-Land Use Cases





# Cassandra-Land Use Cases

- Creating a Keyspace

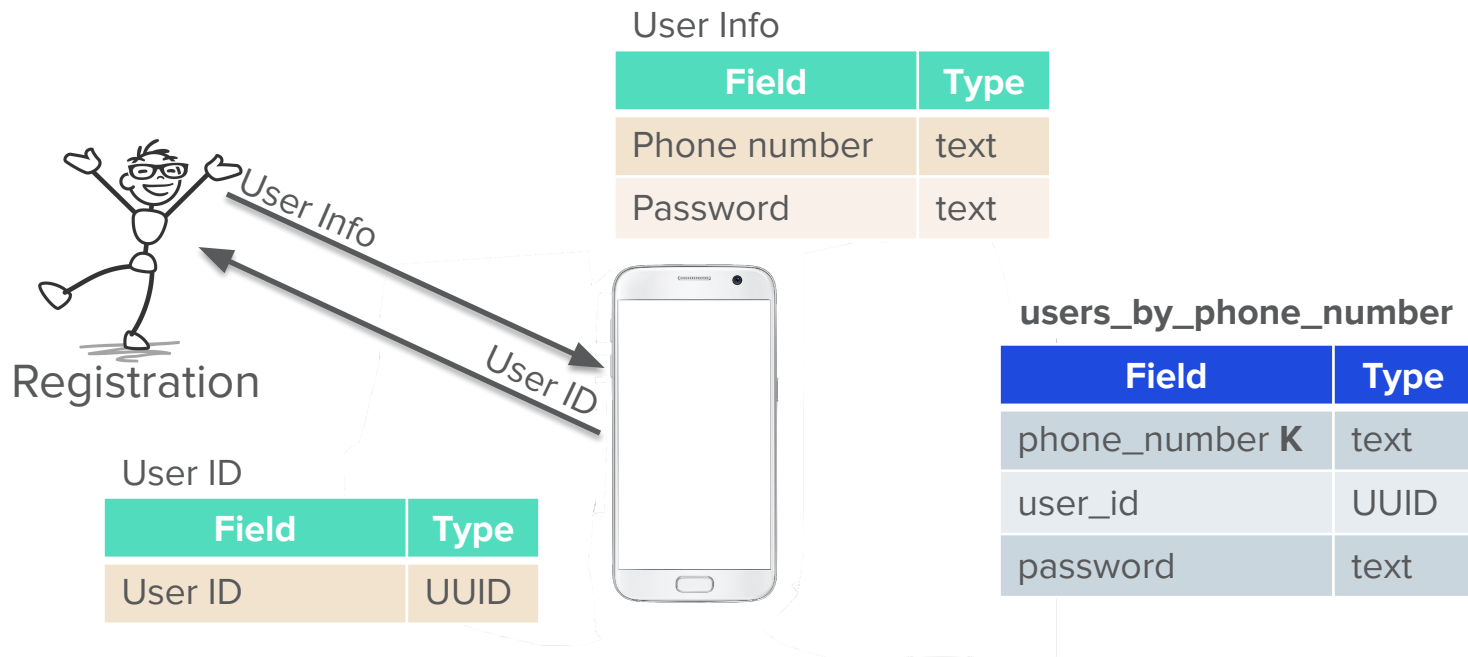
```
CREATE KEYSPACE <keyspace name> WITH REPLICATION = {  
    'class' : <replication strategy>,  
    <datacenter name> : <replication factor>, ... };
```

- For example

```
CREATE KEYSPACE park WITH REPLICATION = {  
    'class' : NetworkTopologyStrategy,  
    'USWestDC': 3, 'USEastDC': 3 };
```



# Cassandra-Land Registration Use Case



# Cassandra-Land Registration Use Case

- Creating a table

```
CREATE TABLE <keyspace name>.<table name> (  
    <field name> <field type>,  
    // Add additional field descriptions here  
    PRIMARY KEY ( <primary key descriptor> )  
);
```



# Cassandra-Land Registration Use Case

- Inserting a row into a table

```
INSERT INTO <keyspace name>.<table name>  
  ( <column list> )  
VALUES ( <column values> );
```



# Cassandra-Land Registration Use Case

- Selecting all rows from a table
  - Typically wouldn't do this in production

```
SELECT * FROM <keyspace name>.<table name>;
```



# Cassandra's Upsert Behavior

- Cassandra does NOT read before writing
- Inserting a row with the same primary key causes an update called an “upsert”
- Similarly, updates to non-existent rows cause an insert
  - Can use a lightweight transaction to prevent an upsert as it does perform a read before writing

```
INSERT INTO keyspace.table IF NOT EXISTS ...
```



# Notebook Data Modeling

---



## Cassandra-Land Project

02-02 - Data  
Modeling:  
Cassandra-Land  
Project PART 1


Developer Day Cluster



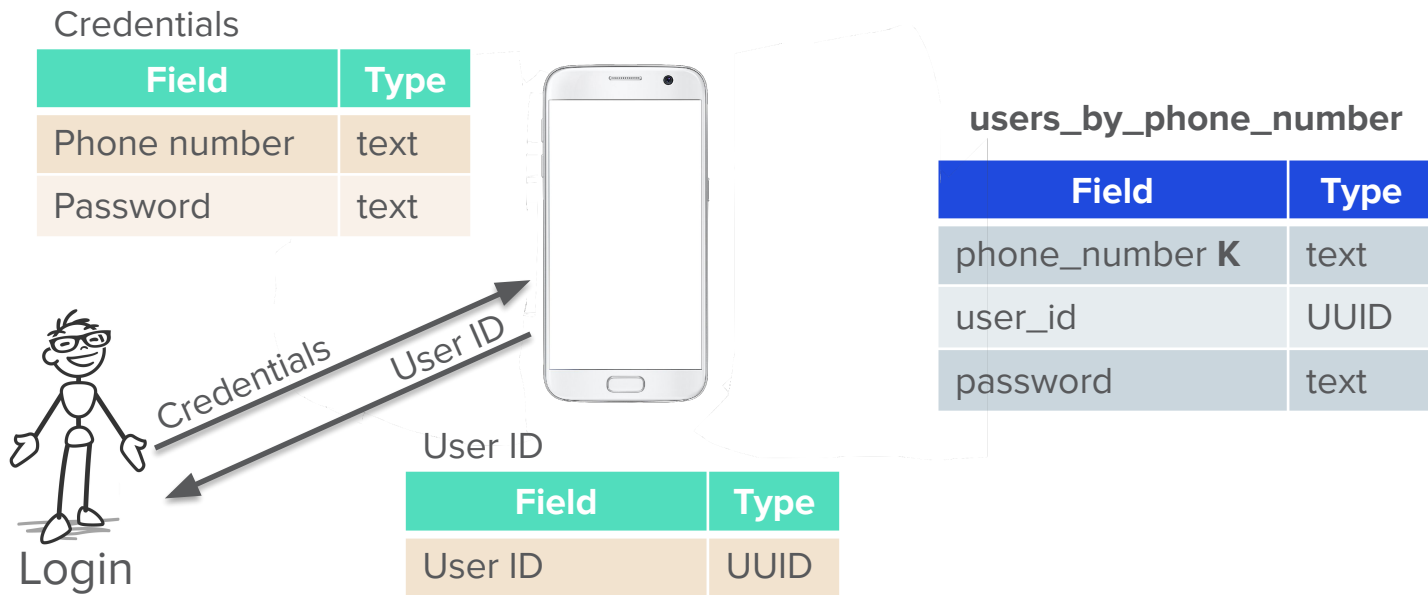
2 days ago



Part 1



# Cassandra-Land Login Use Case





# Cassandra-Land Login Use-Case

- Writing a SELECT statement
  - Must include full partition key
  - Partition keys do NOT support inequalities
  - Not all clustering columns need be specified, but...
  - Any preceding clustering columns MUST be specified

```
SELECT * FROM <keyspace name>.<table name>  
WHERE <query constraints>;
```



# Notebook Data Modeling

---



## Cassandra-Land Project

02-03 - Data  
Modeling:  
Cassandra-Land  
Project PART 2


Developer Day Cluster



2 days ago



Part 2



# Cassandra-Land Ride Alert Use-Case

ride\_instances\_by\_start\_time


Field	Type
start_time <b>K</b>	timestamp
ride_id <b>C↑</b>	UUID
user_id <b>C↑</b>	UUID
ride_name	text
phone_number	text



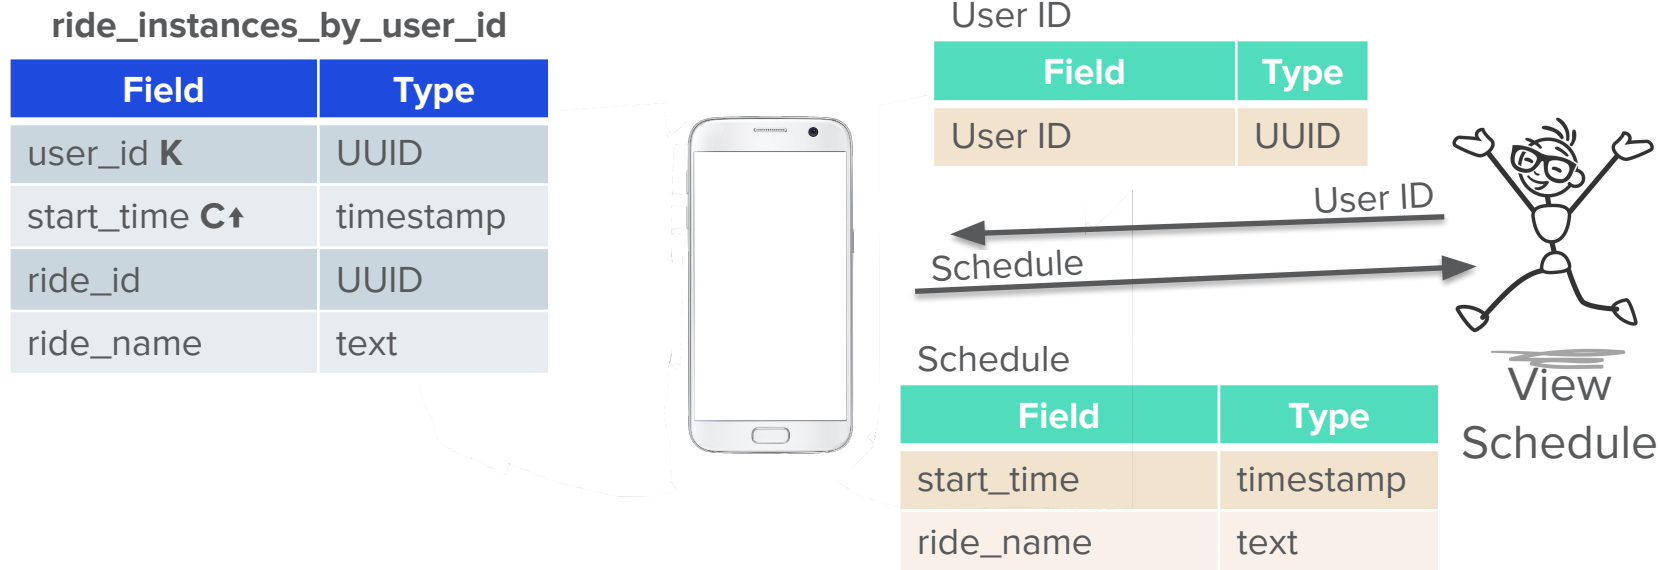
Ride  
Alert

Ride Alert

Field	Type
phone_number	text
ride_name	text
start_time	timestamp



# Cassandra-Land View Schedule Use-Case



# Cassandra-Land Schedule Ride Use-Case

ride\_list\_by\_location

Field	Type
location <b>K</b>	text
ride_id <b>C↑</b>	UUID
ride_name	text
capacity	int

rider\_count\_by\_time\_and\_ride

Field	Type
start_time <b>K</b>	timestamp
ride_id <b>C↑</b>	UUID
rider_count	int

Ride List

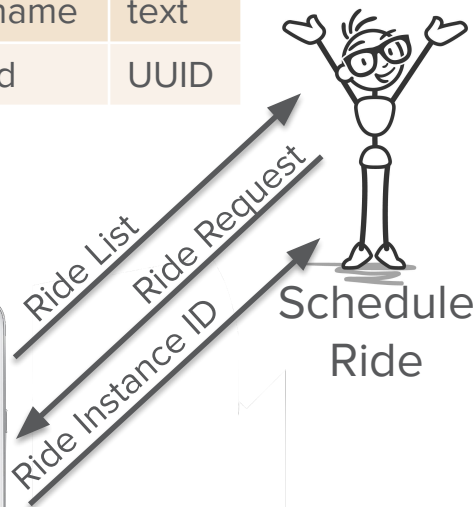
Field	Type
ride_name	text
ride_id	UUID

Ride Request

Field	Type
user_id	UUID
ride_id	UUID
start_time	timestamp

Ride Instance ID

Field	Type
ride_instance_id	UUID



# Cassandra-Land Table Summary

users\_by\_phone\_number

Field	Type
phone_number <b>K</b>	text
user_id	UUID
password	text

rider\_count\_by\_time\_and\_ride

Field	Type
start_time <b>K</b>	timestamp
ride_id <b>C↑</b>	UUID
rider_count	int

ride\_list\_by\_location

Field	Type
location <b>K</b>	text
ride_id <b>C↑</b>	UUID
ride_name	text
capacity	int

ride\_instances\_by\_user\_id

Field	Type
user_id <b>K</b>	UUID
start_time <b>C↑</b>	timestamp
ride_id	UUID
ride_name	text

ride\_instances\_by\_start\_time

Field	Type
start_time <b>K</b>	timestamp
ride_id <b>C↑</b>	UUID
user_id <b>C↑</b>	UUID
ride_name	text
phone_number	text



# Primary Key - What you need to know

- Must have one or more partition key columns
- May have zero or more clustering columns

```
PRIMARY KEY(( <partition key column>,...), <clustering column>,...)
```



# Timestamps

- Format:

```
'YYYY-MM-DDTHH:MM:SS[.fff]'
```

- Notice the quotes
- Milliseconds are optional
- Examples:

```
'2020-01-09T11:45:23'
```

```
'2020-01-09T11:45:23.898'
```





# Update Statement

- Can have multiple <assignment>
- IF is optional – causes a lightweight transaction

```
UPDATE <keyspace name>.<table name>  
    SET <assignment>  
    WHERE <row specification>  
IF <condition>
```



# Batch Statement

- What you need to know – BATCH

```
BEGIN BATCH
    INSERT statement
    INSERT statement
    ...
APPLY BATCH
```

- Once a statement succeeds, Cassandra will ensure all the others succeed
- Can use for inserting into multiple tables



# Notebook Data Modeling

---



## Cassandra-Land Project

02-04 - Data  
Modeling:  
Cassandra-Land  
Project PART 3


Developer Day Cluster



2 days ago



Part 3



# Cassandra-Land

- How to analyze use-cases to derive a data model
- How to denormalize to maintain performance
- How to use lightweight transactions
- How to leverage batch operations





# Thank You

